

SOURAV GHOSH

📧 contactsourav9546@gmail.com 📞 +916203950092 [in](https://www.linkedin.com/in/sourav-ghosh-life/) <https://www.linkedin.com/in/sourav-ghosh-life/> 📍 Pune, India

OBJECTIVE

Experienced Data Engineer with 3+ years of expertise in designing scalable data pipelines, optimizing ETL workflows, and improving database performance to support analytics and business reporting. Proven ability to modernize legacy systems, streamline data integration processes, and implement robust business logic to enhance data accuracy and operational efficiency. Skilled in developing distributed data processing solutions using PySpark with strong focus on data quality through comprehensive testing frameworks. Proficient in advanced SQL, Apache Spark, Python, and cloud platforms such as Azure. Adept at applying machine learning and predictive modeling to help businesses make informed decisions, and leveraging LLMs and AI agents to solve day-to-day tasks and ease developer workloads. Seeking to contribute to a data-driven organization through innovative, high-impact solutions.

EXPERIENCE

Data Engineer

Wipro

📅 July 2021 - Present 📍 Pune, India

Project: Data Migration from Hadoop to Cloud (Databricks & Snowflake)

- Migrated Spark ETL jobs from Hadoop HDFS to Cloud Databricks to Snowflake, setting up infrastructure for Databricks and Snowflake.
- Scheduled ETL workflows using PySpark on Databricks and orchestrated them using open-source Apache Airflow.
- Built scalable ETL pipelines to extract, transform, and load data from various sources like SQL Server, Excel, and flat files into Snowflake using Spark on Databricks.
- Maintained GitHub version control for all transformation logic and notebook updates.
- Developed an ETL pipeline using Databricks notebooks to extract data from relational databases using Spark JDBC connectivity. Utilized Spark for data transformation, including data cleaning, aggregation, and normalization. Loaded the transformed data into Snowflake using the Spark Snowflake connector, enabling efficient querying and reporting.
- Implemented Spark performance tuning techniques like broadcast joins, salting, caching, and persistence, reducing job execution time from 4 hours to 2.4 hours.
- Established secure connections using service accounts to ingest dimension and fact tables from Snowflake to Databricks for final view creation.
- Ensured high-quality documentation of workflows, enabling smoother developer onboarding.
- Developed and implemented automated ETL pipelines using Azure Data Factory to extract, transform, and load data from various sources into Azure Blob Storage, ensuring seamless integration with Azure Databricks for advanced data processing and analytics.
- Built and deployed scalable data processing workflows using Azure Data Factory and Azure Databricks, utilizing Spark clusters for data transformation, aggregation, and storage in Azure Blob Storage, optimizing data handling performance.
- Enabled CI/CD integration with GitHub for version control, automated deployment, and collaborative development, enhancing code quality, traceability, and rollback capabilities.
- Managed ingestion of raw data into Amazon S3 and leveraged AWS Glue to automate ETL processes for data cleansing and transformation.

Data Engineering Project

- Optimized Azure Data Factory (ADF) pipelines to ingest and process various healthcare claim data from sources such as REST APIs, JSON files, SQL databases, and CSV files. Achieved an improvement 20% in data integration efficiency, significantly reducing data processing and reporting times.
- Architected a scalable Medallion architecture (Landing → Bronze → Silver → Gold) on Azure Data Lake Storage Gen2, ensuring organized, traceable data processing layers from raw ingestion to refined business-ready datasets.
- Supported and enhanced Master Data Management (MDM) processes and data quality checks.
- Orchestrated large-scale ETL workflows in Azure Data Factory (ADF) and Databricks to support downstream analytics.
- Executed complex transformations including joins, aggregations, and filters to meet business logic requirements.
- Designed star schemas and implemented dimensional data modeling to optimize reporting and analytics.
- Successfully migrated SharePoint applications and databases from one on-premises server to another, ensuring seamless data integrity and minimal downtime.
- Designed and developed a Power BI dashboard report displaying Service Delivery counts for CARS and Incidents, enabling data-driven insights to maximize productivity.
- Supported application deployments through Azure DevOps during overnight schedules, ensuring successful and uninterrupted release cycles.
- Developed a monitoring script to check the availability of .NET web applications and automatically trigger email notifications to the respective stakeholders for follow-up.
- Created comprehensive SOPs and detailed technical documentation to streamline the onboarding process for new team members.
- Optimized slow-performing OutSystems pages and APIs by identifying heavy queries, unoptimized aggregates, and server bottlenecks, significantly improving application response times.
- Performed SSL certificate renewal and installation for OutSystems environments, ensuring uninterrupted secure communication for applications and platform services.
- Successfully upgraded the OutSystems Platform Server to version 11.18.1 and the OutSystems Lifetime Management Console to 11.16.2 across Test and Production environments, leading efforts to troubleshoot and resolve upgrade-related issues to ensure smooth deployments.
- Conducted backup and recovery operations, ensuring minimal data loss and high availability of SharePoint content.
- Automated administrative tasks using PowerShell scripts for SharePoint and Microsoft 365 environments.
- Managed the end-to-end deployment process for Windows SharePoint Solution Packages (WSP) across multiple SharePoint environments (development, testing, and production).
- Delivered 30+ SharePoint installations and application configurations, covering farm setups, service applications, and web application deployments across multiple environments.
- Designed and implemented a metadata-driven ETL pipeline using Azure Data Factory (ADF), enabling automated, dynamic ingestion of structured and semi-structured data (SQL, flat files, and API data) through configuration-based orchestration.

SKILLS

PythonPySparkSQLSpark

DatabricksAzure SynapseHadoop

SnowflakeData modelingDatabase

DatawarehouseETLAzure

Azure Data FactoryData Pipeline

CI/CD JenkinsPower BIGit

Delta LakeHDFS

EDUCATION

Bachelors of Computer Application

Arka Jain University, Jamshedpur

📅 2018-2021

CGPA : 8.0

Senior Secondary School

Kerala Public School,Kadma

📅 2018

PERCENTAGE : 66%

10th Standard

Kerala Public School,Kadma

📅 2017

Percentage : 78.0%

CERTIFICATION & REWARDS

- AWS Certified Cloud Practitioner
- Machine Learning Specialization-Andrew Ng
- SEP Rising Star Engineer Q3-2024.
- Q2-2025 Catalyst Award Winner for the impactful contributions over and above the BAU
- Runner up in AWS DeepRacer - Mumbai Region
- Certification of Achievement in SQL Bootcamp by Great Learning
- Certification of Achievement in Snowflake by Snowflake
- Certification of Achievement in Big Data by Databricks

PROJECTS

Olympic Data Analytics Pipeline Using Azure Data Factory, Databricks Synapse

Self Project

📅 October 2025 📍 Pune

- Designed and implemented a scalable data pipeline using Azure Data Factory to ingest Olympic datasets into Azure Data Lake Storage Gen2, organizing data in Bronze, Silver, and Gold layers for efficient processing.
- Developed data transformation workflows in Azure Databricks using PySpark to clean, enrich, and aggregate athlete performance and medal data, enabling country-wise and sport-wise insights.
- Loaded curated data into Azure Synapse Analytics and integrated Power BI dashboards to visualize key Olympic metrics such as top-performing countries, athlete counts, and event-wise medal distributions.